

ADVANCED C WORKSHOP

Day 4
Overview of the Computer Architecture





- How the CPU does what it does
 - Understand how the CPU in your computer works
 - Build a good model of your CPU in your mind
- How can we use that to write better code
 - Focus on “hot” code
 - Use matrix multiplication as an example
- Focus on x86_64, Linux



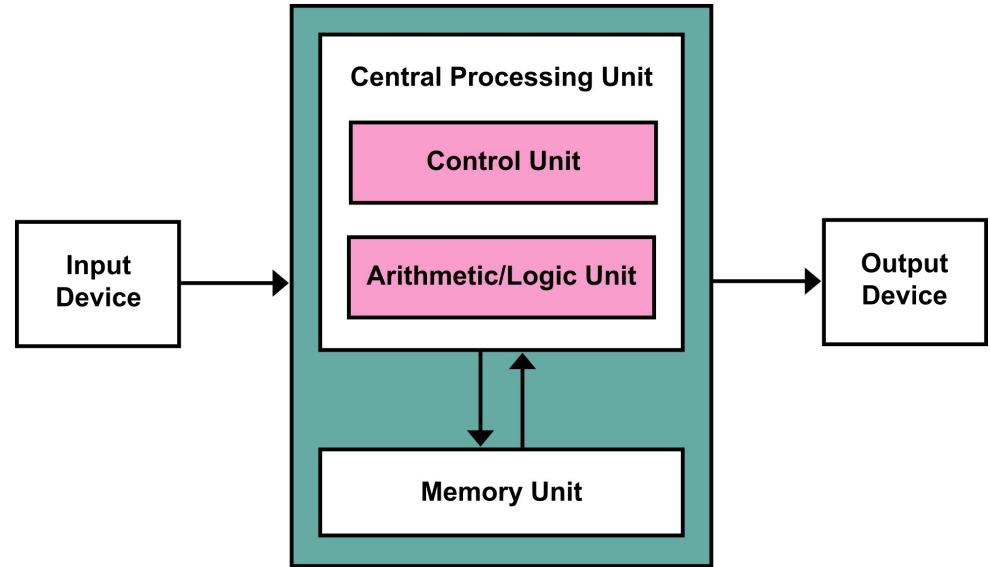
Why study architecture

- It is cool and fun!
 - Solve real world problems
 - Solid engineering solutions
 - Very active area of research
- It will help us write better and faster code
 - Moore's Law and friends are maybe dead (kind of)
 - Atoms are $\sim 0.5\text{nm}$, transistors $< 50\text{nm}$



What the CPU does

- Crunch numbers
- Input, Process, Output



Von Neumann Architecture



Different kinds

- Registers - what CPU needs right now
- RAM - everything else needed to do a job
- Secondary storage - things we might need later



- Small but very fast storage (~1 cycles)
- Built right in the CPU
- Can be general purpose
 - x86_64 has 16 general purpose 64-bit registers
 - Arm (and other RISCs) have >32
- Can be specialized
 - Store state (PC, SP, offsets etc)
 - SIMD etc.



RAM

- Larger memory, decently fast (~100 cycles)
- Runs independently from the CPU
- Different clock speed



- Makes our life easier
- Idealized abstraction over storage resources
- Illusion
 - RAM is very large
 - Your process is the only thing running
 - Memory is linear and unfragmented
 - You ask for 4GB, you get nice clean 4GB block



- Makes our life much easier
- We can move pages between RAM and disk
 - Linux has swaps, Windows has ReadyBoost
 - You can run 10 processes that each need 1 GB even if you only have 4GBs of RAM
 - Performance hit is substantial, but it's better than giving up



- Files and I/O devices can be mapped to RAM
 - Use the RAM as a buffer
 - Faster reads and writes
- Programs can be memory position independent
 - Useful for dynamic/shared libraries (plugins)
- Trick OS into using more memory than there is

1[|||||] 6.6% 800MHz 48°C]
2[|||||] 6.1% 800MHz 49°C]
3[|||||||] 11.3% 838MHz 48°C]
4[|||||] 4.6% 800MHz 46°C]
Mem[|||||||||||||||||] 3.63G/15.3G]
Swp[0K/16.0G]
Network: rx: 0KiB/s tx: 0KiB/s (1/1 pkts/s)

PID	USER	PRI	NI	VIRT	RES	SHR	S	CPU%	MEM%▽	TIME+	Command
4145	abhigyan	20	0	13.2G	679M	280M	S	2.0	4.3	45:46.43	firefox
17581	abhigyan	20	0	3178M	443M	121M	S	11.2	2.8	4:36.87	firefox -contentproc
13918	abhigyan	20	0	2469M	311M	184M	S	2.6	2.0	0:31.61	nautilus
4462	abhigyan	20	0	3124M	290M	101M	S	0.7	1.9	0:29.26	firefox -contentproc
4487	abhigyan	20	0	2692M	222M	99M	S	0.0	1.4	0:06.18	firefox -contentproc
3959	abhigyan	20	0	852M	213M	58684	R	11.9	1.4	3:51.12	emacs --daemon
13248	abhigyan	20	0	2608M	209M	98M	S	0.0	1.3	0:51.03	firefox -contentproc
4857	abhigyan	20	0	27.4G	201M	91220	S	0.7	1.3	1:38.56	firefox -contentproc
20524	abhigyan	20	0	2631M	184M	97M	S	0.0	1.2	0:07.70	firefox -contentproc
3660	root	20	0	26.4G	183M	101M	S	7.3	1.2	17:30.49	Xorg :0 -seat seat0
4386	abhigyan	20	0	2553M	179M	95268	S	0.0	1.1	0:08.47	firefox -contentproc
4565	abhigyan	20	0	2522M	172M	98764	S	0.0	1.1	0:08.36	firefox -contentproc
19975	abhigyan	20	0	2507M	160M	97728	S	0.0	1.0	0:03.26	firefox -contentproc
21240	abhigyan	20	0	1365M	159M	134M	S	0.0	1.0	0:00.28	flameshot
20698	abhigyan	20	0	2478M	159M	95732	S	0.0	1.0	0:02.02	firefox -contentproc
20046	abhigyan	20	0	2484M	155M	94888	S	0.0	1.0	0:03.07	firefox -contentproc
20719	abhigyan	20	0	1342M	154M	122M	S	0.7	1.0	0:01.79	alacritty
11794	abhigyan	20	0	2525M	153M	97464	S	0.0	1.0	0:10.19	firefox -contentproc
20368	abhigyan	20	0	2479M	151M	96376	S	0.0	1.0	0:02.11	firefox -contentproc
4341	abhigyan	20	0	2582M	150M	99124	S	0.0	1.0	0:08.81	firefox -contentproc
4441	abhigyan	20	0	2505M	149M	93936	S	0.0	1.0	0:04.87	firefox -contentproc
20318	abhigyan	20	0	2479M	149M	91712	S	0.0	1.0	0:01.30	firefox -contentproc
4337	abhigyan	20	0	2946M	149M	101M	R	6.0	1.0	11:40.45	firefox -contentproc
11734	abhigyan	20	0	2527M	148M	97756	S	0.0	1.0	0:07.58	firefox -contentproc

F1Help F2Setup F3Search F4Filter F5Tree F6SortBy F7Nice -F8Nice +F9Kill F10Quit



- Address translation has overheads
- MMU does this to free the CPU
- MMU is a hardware unit in the CPU
- Kernel populates tables in the MMU
- MMU translates on the fly
 - We can cache!




- Caches virtual address to physical address translation
- Is usually hierarchical
- May have separate sections for small and large pages
- Typical TLBs have ~1000 entries, <1% miss rate
- Miss penalty is ~100 cycles
- TLB thrashing - lots of TLB misses





Instruction Cycle

- Cycle that a CPU follows until powered off
- Fetch - read next instruction from RAM[PC]
- Decode - Interpret the instruction
- Execute - Do the actual work
- Write - store the results
- Repeat





Pipelining - Instruction Level Parallelism





- Sometimes we don't know which instruction will come next
- Conditional Jumps / Branches
- Wait until we know exactly what will happen?
 - Bad, we have to flush the pipeline and start over
- Choose one randomly?
 - Great only when we are right



- Use some strategy to predict outcomes
- Simple strategies work well
 - Branches are always taken
 - Backward branches are taken, forward branches are not
- We can go further
 - Cache the paths taken previously to help BP
 - Use Neural Networks (because why not?)



- We cannot override BP, but we can help
- Modern processors guess that backward branches are taken
- We can ask the compiler to lay code out to help with this

```
if (__builtin_expect(foo==bar)) {...}
```



- Each trip to the RAM takes ~ 100 cycles
- For context, integer addition is ~ 1 cycle (pipelined)
- Let's cache data that we might need
 - If you access x , you might access x 's neighbors too
 - Spatial locality
 - If you access x 10 times, you might access it again
 - Temporal locality




- Practical CPUs have multiple levels of cache
 - L1 cache - 32KB, 5 cycles
 - L2 cache - 1024KB, 70 cycles
 - L3 cache - 10MB, 80 cycles



- Not everything fits in the cache
- To add, we need to evict
- Algorithms choose pages to evict
 - Least Frequently Used
 - Least Recently Used
- L1 cache has simple eviction strategy
- Higher cache levels are more complicated



Memory Hierarchy





Same Instruction, Multiple Data

- Find the sum of a large array (databases)
- Multiply a lot of number pairs (matrix multiplication)
- Invert, XOR a large stream of data (hash functions)



- Modern processors have SIMD instructions
- Intel and AMD have SSE, AVX1, AVX2 etc
- Arm has Neon
- SIMD enables data parallelism
 - You can do the same thing on multiple pieces of data
- SIMD is not concurrency
 - You cannot do multiple kinds of things at the same time



- CPU has different circuits and units for different things
- Idea: do multiple things at once
- Only possible with careful data dependency checks
- Enables execution of multiple instructions per clock cycle



Simultaneous Multi-Threading (SMT)


- Multiple execution units in the same CPU
- With careful pipelining, CPUs can ~double the performance
- Eg: Hyperthreading, AMD SMT
- (X cores, $2 \times X$ threads)



- Everything so far happens in one core
- Most CPUs have multiple cores
- Some problems are embarrassingly parallel
 - Use N CPUs, get close to $N\times$ boost



- Do work before you need to do it
- Needs careful planning and dependency checks
- Eg: Branch Prediction
- Thread Level Speculation - SE in multithreading systems
- Has caused problems in the past
 - Look into Spectre, Meltdown etc



ADVANCED C WORKSHOP

THANK YOU!

