



Multithreading - I

Ashutosh Bohara




Outline

- Threads
- Atomic Operations
- Volatile variables
- Mutex
- Semaphores

Scheduling

- Selecting a “job” for the CPU to perform
- Performed by the OS
- Several popular strategies:
 - ◆ CFS (Linux)
 - ◆ Round Robin
 - ◆ FCFS



Context Switching

- Performed by the scheduler(a part of the OS) when changing the thread currently running in a CPU core
- Saves all the relevant information like program counter, registers, page tables, etc
- Cost of context switching






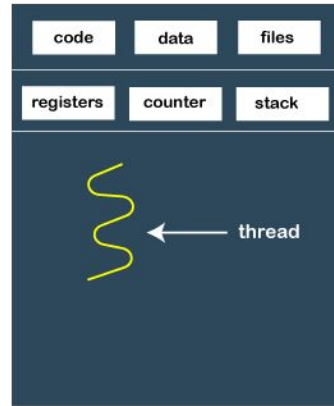
Why Multithreading ?

- A CPU has many cores
- Why not use all of them?
- Parallelizable tasks
 - ◆ Matrix manipulation (used in games, ML)
 - ◆ Image processing
 - ◆ Blockchain
 - ◆ Servers


True Parallelism vs Task Interleaving



Processes and Threads



Single-threaded process



Multi-threaded process



Thread

- Single execution unit that can get scheduled by the OS to run on a CPU core
- Single sequential flow of control
- A process can create many threads, all of which share its virtual address space and system resources



Hardware threads vs Software threads

- CPU with many cores and threads
- One hardware thread per CPU core (with the exception of hyperthreading and SMT)
- Software threads are a virtualization over the hardware threads
- When there are more software threads than hardware threads, at least one of the software threads is always sleeping

From here onwards “thread” will mean a software thread



Thread Operations

- Creating threads
- Joining threads
- Terminating threads
- Detaching threads
- Sleep



Let's create a thread



Infinite threads, Infinite Power ??




Thread Support in C

- C11 added support for multithreading with `threads.h`
- Since it was termed optional, MS still hasn't added support for the library
- Introducing “`threading.h`”



Passing Arguments

- Provided during thread creation
- `void*` is used to pass arguments of arbitrary type to the function
- Arguments passed to the child thread must be valid while it is being used in the child thread (code sample)



Adding upto 100000 using multiple threads



Data race/ race condition

- Conflicting accesses to the same memory location from multiple threads
- In the previous program, the incrementation operation creates conditions for data race.

Ideal Case

Thread 1	Thread 2		Integer value
			0
read value		←	0
increase value			0
write back		→	1
	read value	←	1
	increase value		1
	write back	→	2

Without Synchronization


Thread 1	Thread 2		Integer value
			0
read value		←	0
	read value	←	0
increase value			0
	increase value		0
write back		→	1
	write back	→	1



Atomic Operations


- CPU operations that cannot be broken down into smaller steps having an observable effect
- Once started, an atomic operation cannot be left incomplete
- x86 has 'lock' prefix added before certain instructions to make them atomic.
 - ◆ e.g. *lock add*, *lock inc*, *lock sub*

inc



The effect is
visible globally

atomic inc



Performed all at once



Compiler Ininsics

- Atomic Operations are available in C through intrinsics, which are compiler specific
- But 'threading.h' wraps these into a common interface
 - ◆ InterlockedIncrement
 - ◆ InterlockedDecrement
 - ◆ InterlockedAdd
 - ◆ InterlockedCompareExchange



Let's add Atomic Increment to our program



Compiler Optimizations

- Compilers are very sophisticated and impressive programs
- What ?, not How ?
- Some of the basic optimizations performed by a compiler are
 - ◆ Loop unrolling
 - ◆ Compile time calculation
 - ◆ Variable caching
 - ◆ Function inlining
 - ◆ Dead code elimination
 - ◆ Code reordering

```

int foo(){
    int res = 0;
    for(int i = 0; i <= 10000; ++i) {
        res += i;
    }
    return res;
}

```

```

1  _foo    PROC
2          xor     eax, eax
3          mov     ecx, eax
4  $LL4@foo:
5          add     eax, ecx
6          inc     ecx
7          cmp     ecx, 10000
8          jle     SHORT $LL4@foo
9          ret     0
10 _foo    ENDP

```

```

1  foo:
2      mov     eax, 50005000
3      ret

```



Variable Caching ?? Seems problematic!



Volatile

- Type qualifier keyword in C
- Information for the compiler that its value may change at any time.
- It is widely used in embedded programming where a register that can be modified by an I/O operation is directly mapped to a variable
- Value of a variable may be changed by some other thread



Let's add Volatile to our program

Questions?



Memory Barriers

- Compiler can reorder segments of code
- But that's a problem for multithreaded programs (let's see how)
- A *software* memory barrier tells the compiler the compiler to not reorder memory operations across the barrier.





Hardware Barriers

- The CPU can also reorder instructions (out-of-order processing)
- We need a **hardware** memory barrier that prevents the reordering of CPU instructions.
- Unlike software barriers, hardware barriers generate actual assembly code
 - ◆ lfence, mfence, sfence ,
- Compiler specific, available as intrinsics



Limitations of Atomic Operations

- There are only so many instructions a CPU can perform atomically
- Writing to a file from multiple threads?
- Read/write from/to a console from different threads?
- “atomic like” behavior, but more generalized



Mutex

- Mutex (Mutual Exclusion) is a simple data structures with states:
 - ◆ Locked
 - ◆ Unlocked
- Threads that share data use mutex to control access to it
- Critical section is executed only if mutex is not locked by another thread by locking the mutex first



Critical Section

- Must execute in isolation from rest of the program.
- Should be no interference from other threads
- Access to the shared data.



Mutex Example



Mutex Implementation

```
typedef struct Mutex {  
    uint32_t value;  
} Mutex;
```



Mutex Implementation

```
static void MutexInit(Mutex *mutex) {  
    mutex->value = 0;  
}  
  
static void MutexLock(Mutex *mutex) {  
    while(mutex->value) {}  
    mutex->value = 1;  
}  
  
static void MutexUnlock(Mutex *mutex) {  
    mutex->value = 0;  
}
```

What happened ??

```
MutexInit PROC                                ; COMDAT
    mov     DWORD PTR [rcx], 0
    ret     0
MutexInit ENDP

mutex$ = 8
MutexLock PROC                                ; COMDAT
    mov     DWORD PTR [rcx], 1
    ret     0
MutexLock ENDP

mutex$ = 8
MutexUnlock PROC                              ; COMDAT
    mov     DWORD PTR [rcx], 0
    ret     0
MutexUnlock ENDP
```



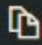
How do we fix that ??



What now ??

- The problem is same as before.
- The access to the variable value is not atomic.
- **How do we make it atomic ??**

C++

 Copy

```
LONG InterlockedCompareExchange(  
    [in, out] LONG volatile *Destination,  
    [in]      LONG          ExChange,  
    [in]      LONG          Comperand  
);
```


Questions?



Semaphores

- Generalised mutex
- Used to manage concurrent processes by using a simple integer value
- An integer variable that is shared between threads
- Binary semaphores vs counting semaphores

Mutex vs Binary Semaphore ??





Mutex Ownership

- Mutexes have a concept of ownership
- Unlocking a mutex by someone not its owner is undefined
- But semaphores are used to do exactly that
- A semaphore locked by a worker thread will be unlocked by the work creating thread

Let's look at an example



Deadlocks

- When two or more threads are waiting for each other
- Two mutexes waiting for each other to release its lock
- When a thread is deadlocked, it can not run



Avoiding Deadlocks

- Timeouts
- Changing locking order
- Synchronisation techniques / Condition Variables
- Lock-free programming

Thank You !!
